Persistent K-Coverage UAV–IoT: Complete Project Context

Professor progress meeting — 4 November 2026

Himanshu Nagwanshi, Surya Prakash S, Harsh Varshney, Gaddam Viswavijith

Agenda (15–20 minutes)

1. The problem, underspecification, and what the project actually is
2. What changed: old midsem report $\rightarrow$ current report
3. The current report: the artifact and the evaluation protocol
4. The measured results: where learning wins, where it loses
5. The feasibility certificate and energy ceiling
6. Planning, distillation, classical baseline
7. The audit and the OMNeT/INET co-simulation
8. What is real, what is not real, and what we need from you

Executive summary

What this project is

- A *feasibility-certified evaluation study* of persistent K-coverage UAV-IoT over IoT sensors, with decentralized MARL control.
- It does *not* claim RL is universally superior. It measures when learning wins, when it loses, and when the instance is infeasible.
- The artifact is a 27-page report plus an executable INET co-simulation and claim audit.

Where it stands

- Corrected MAPPO beats the strongest rule by $+0.028$ on nominal (0.9512 vs 0.923).
- Learning wins on low connectivity ($+0.125$) and both failure regimes; loses on small battery and spread start.
- Config is *not* persistently energy-feasible (ceiling 0.113 at steady state; 400 s is a finite sprint).

The actual problem

Scenario

- $N = 8$ active UAVs + 1 reserve
- $M = 20$ Gaussian sensors in an 80×80 m field
- $K = 2$ required coverage; connectivity must stay intact
- 2 charging pads; clustered corner start; 400 s missions
- Actions: hover / 4 moves / charge
- Observation: local + communication neighbourhood + 4×4 deficit map + ego bearings

Why it is hard

- Persistent, eventually infeasible: battery and pads cannot keep 8 UAVs up indefinitely
- Coverage and connectivity conflict; the swarm must stay connected while deploying across the field from one corner
- The deployed controller has no global state; it sees local observations only

What you are shown today

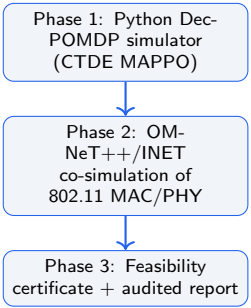
- A 27-page current report with an achievability certificate
- An OMNeT++/INET 4.6 co-simulation at 802.11 PHY/MAC level
- An audit of every claim in the report

The configuration we test

Parameter	Value
UAVs (active + reserve)	$8 + 1$
Sensors M	20
Coverage target K	2
Region Ω	80×80 m
Coverage radius R_{cov}	35 m
Communication radius R_{comm}	40 m
Altitude z_0	50 m
Step Δt	0.2 s, $v_{\max} = 5$ m/s
Safety distance	5 m
Battery $B^{\max}$	30,000 J
Hover / move / avionics	50 W / 75 W / 5 W
Charge rate	0.0002 of pack per step
Episode	400 s (2000 steps)
Deployment	connected cluster in a corner

Initial deployment covers only 3 of 20 sensors; a centred ring would cover 17 and make the comparison degenerate ($CK = 0.850$).

The two Halves of the System



- INET is the fidelity check, not a claim that the two simulators are identical.
- The report's claims audit recomputes exact quantities; it does NOT audit experimental CK tables.
- The INET example now runs headless and records scalars to `General-#0.sca` at $t = 40$ s, event #71885.

Old midsem report → current report

Dimension	Old report	Current report
Length	10 pages; claims deferred	27 pages; measured claims with intervals
Framing	MARL framework proposal	Feasibility-certified evaluation study
Can it be measured?	Not cleanly; no frozen policy	Frozen policy; held-out scenes; paired CIs
Headline	None verified	Learner 0.9512 vs rule 0.923; classical scheduler 0.929
Where it wins	–	Low connectivity
Where it loses	–	Small battery; spread start
Persistence	Assumed	Not persistently feasible; 400 s sprint only
Metrics	SC / utilisation-driven	Non-degenerate metrics; refresh time; connectivity
INET evidence	None	802.11 MAC/PHY; BonnMotion; $t = 40$ s; event #71885
Honesty	Not reported	Losses reported; appendix; 38 audit items, 0 failing

The comparison is against the original 10-page midsem submission, not a hypothetical ideal.

The certificate (Section III-G) – what it actually says

- **Region bound:** $N \geq 4K$ for the square, because corner discs are pairwise disjoint at 80 m vs $2R_{\text{cov}} = 70$ m. Tight for this config.
- **Packing bound:** for a given sensor layout, $N \geq K \cdot P(A, R_{\text{cov}})$. Excludes $N = 5$ on a measured 13% of random draws.
- **Oracle ceiling:** teleporting placement oracle upper-bounds any policy. At the reported configuration the ceiling is 1.000, i.e. coverable.
- **Metric:** the all-conditions-simultaneously service-continuity metric is identically zero for every policy in the infeasible set; we replace it with graded metrics (time-averaged deficit, max refresh time).
- **Verified:** 38 items recompute; 0 failing; 170/170 determinism; corner construction is admissible and connected.

The energy-flow certificate (Section III-H) – the finding we did not expect

The arithmetic

- 2 pads $\times$ 24.9 W net $=$ 49.8 W returned
- 1 airborne UAV draws 55.1 W
- Max sustainable airborne fleet $=$ 0.90 UAVs
- $K = 2$ demands 8 UAVs $\Rightarrow$ persistent ceiling 0.113

What that means

- At $T = 400$ s: ceiling 1.000 (only by draining initial packs)
- At $T = 2000$ s: ceiling 0.367 (full) / 0.326 (mean 0.85 charge)
- Only $N = 12$, 4 pads, 60 kJ, 0.0004 charge rate reaches 1.000
- The reported configuration is **not persistently feasible**; it is a 400 s broadcast, not a long-term mission

The actual ML machinery (Section IV)

- **Algorithm:** parameter-shared CTDE MAPPO; centralized critic during training; decentralized actors at deployment
- **Observation:** 68 dims — 30 base (position/velocity/battery/mode/neighbourhood) + 16 map grid + 8 agent ID + 3 pad bearing + 11 ego-centric bearings to $k = 3$ nearest deficits + neighbour repulsion
- **Action spaces:** 6 discrete actions (hover, 4 moves, charge); *all* actions masked by safety filter; articulation-point exclusion
- **Policy step:** sample during training; deployment follows the sampled distribution (not argmax)
- **Failure recovery:** a dedicated reserve UAV re-seats a failed slot at a pad
- **Charging:** pad allocation/arbitration; pad selection handled by the scheduler
- **PHY gate:** 802.11-derived link budget decides whether a UAV can cover/neighbour; the simulator and INET are separately verified

The measured headline (nominal, corrected protocol)

Arm	CK	Interval
Strongest hand rule (Reactive-local)	0.923	–
MAPPO ego + shaping	0.9512	[0.938, 0.964]
Classical OR scheduler	0.929	20 scenes, $\sigma = 0.062$
Oracle ceiling	1.000	deterministic (every layout)

- MAPPO beats the strongest rule by $+0.028$, interval excludes 0
- Gain concentrates in low connectivity; not in energy or start regimes
- The classical OR scheduler still trails the learner by 2.2 percentage points

Stress regimes (frozen policy; zero-shot)

Regime	Rule	MAPPO	Classical	MAPPO − rule
Nominal	0.921	0.947	0.929	+0.026
Small battery	0.829	0.765	0.791	−0.064
Low connectivity	0.804	0.928	0.808	+0.125
Spread/start	0.997	0.959	0.980	−0.038
Failure	0.912	0.946	0.919	+0.034
Multi-failure	0.910	0.945	0.919	+0.035

Frozen policy, five held-out scenes per regime, same physical setup; single_pad changes the architecture so no zero-shot policy can load.

Classical = oracle formation + Hungarian assignment + rotation (not a weak hand-tuned heuristic).

Planning oracle vs the deployable CTDE policy

- **Centralised planning** (H=2, K=4, 500 steps, 20 scenes): 0.9531 (200 step) / 0.9723 (2000 steps); paired +0.0211, CI [+0.014, +0.028], resolved
- It rolls the *global* simulator and picks a joint action — not deployable by the decentralized controller
- **Distilled 20-epoch BC actor**: 0.9618 vs base 0.9512; paired +0.0106, CI [-0.002, +0.023] (directional, not resolved at $n = 5$)
- Distillation moves about half of the oracle's gain into the local actor without global state
- **Classical OR scheduler:** 0.929 nominal; loses to MAPPO on every same-architecture regime; cannot fix small battery either

INET co-simulation: what it verifies and what it does not

INET is execution support, not a replacement for the Python simulator.

- INET 4.6 / OMNeT++ 6.4; 8 UAVs, 20 sensors, BonnMotion trace; 40 s run; event #71885; End.
- 802.11 MAC/PHY path loss (`BreakpointPathLoss`) vs Python reference: algebraically identical to 0.012 dB over 24 points
- 170/170 determinism check on the same seeds
- **What it does not show**: persistent operation (battery drains in 400 s); hardware calibration; a deployment guarantee
- Network margin: at 80 m scale, range 1762 m $\Rightarrow$ geometry/battery/coverage bind first, not the radio

Claim audit and what it does / does not cover

The audit

- 38 items; 34 pass; 4 informational; 0 fail
- Verifies: corner packing bound, oracle construction, device sensitivity, determinism (170/170), PHY equivalence, phase-2 execution
- Saves `claim_audit.json`; re-run produced identical verdicts

Out of scope

- It does *not* audit Table VIII coverage
- It does *not* choose the safest scenario; all numbers trace to artefacts but not peer review
- It does *not* certify the deployment, only the model

What the project actually contributes

1. **A feasibility certificate** that stops you from reporting an unachievable number as a controller result
2. **A corrected evaluation protocol** that freezes the policy and benches it on held-out scenes with paired intervals
3. **Honest negatives** — where learning loses (small battery, spread start) and why the gap exists
4. **Energy-flow ceiling** — the certificate that says the current benchmark is a 400 s sprint, not persistence
5. **One honest positive** — low-connectivity regime where learning wins clearly and consistently

What you should not take away from the talk

- Not "RL beats every heuristic" — it loses 2 of 6 regimes
- Not "persistent" — the config depletes in 400 s; the mission is not indefinite
- Not a new RL algorithm — it is a study of *when the controller works*
- Not hardware-validated — it is INET + Python, on a fixed CPU profile
- Not a top-tier paper — the contribution is discipline, not method novelty

Open issues (none hidden)

1. Observation documentation is wrong inside the report: the evaluated model uses map-grid + agent-ID + pad bearing (68 dims total); the text still implies they were excluded
2. The provenance of the 0.862 "full coverage state" result is not audited in the current code path (it measured 57-dim)
3. "Persistent" is a misnomer at 400 s; title and framing overstate the horizon
4. The two contributions are repackaged known results (packing duality; conservation of energy)
5. No hardware validation; no real wind / battery ageing; no multi-hop relay in dynamic topology

What I need from you, sir

1. **Scope decision:** (a) keep this as a conference submission with the honest framing; (b) add the two missing experiments (proactive staggered recharge; R_{comm} sweep) before submission; (c) drop the submission entirely
2. **Title/abstract decision:** keep energy-first wording, or move the certificate into the abstract and lead with the system
3. **Proposed next work:** (i) fix documentation mismatch; (ii) audit 0.862 provenance; (iii) R_{comm} sweep with the frozen policy; (iv) if results justify, run proactive recharging variant
4. **Risk acceptance:** the contribution is honest but marginal; the professor's name is on the report only if you approve

Next two weeks — the concrete plan

1. Week 1: fix the observation documentation inside the report; audit every claim in Table VIII against artefact; recompile and re-run the claim audit
2. Week 1: $R_{\mathrm{comm}} \in \{22, 26, 30, 34, 38, 42\}$ sweep with frozen policy, Reactive-local rule, classical scheduler; record CK and mask-blocked fraction
3. Week 2: implement proactive staggered recharge (early recharge when pad free and $b_i < 0.70$); retrain only if you approve it as the headline contribution
4. Week 2: final report package (27-30 pages, appendix with the measurement correction record, references current, claim audit clean)

Demo pointer (no fabrication; all reproducible)

- **OmerNet run:** `wsl -d Ubuntu -- opp_env run inet-4.6.0 -c 'cd .../uav_omnet && opp_run -u Cmdenv ...'` → event #71885, $t = 40$ s, End.
- **Claim audit:** `python ml_impl/analysis/claim_audit.py` → 34 pass, 0 fail
- **PPT:** `PPT_SI_CHANGES.pdf` — this deck
- **Report:** `Mlawn-midsem-tex.pdf` — 27 pages, current title
- **Context bundle:** `PROJECT_REPORT_CONTEXT.txt` — for independent review

Three bottom lines

1. **What we have:** a correct, reproducible, honest study that shows where learning wins and loses, why the configuration cannot persist, and why the metric was misleading
2. **What we do not have:** a novel RL algorithm, a persistent mission, or a hardware result — and the paper does not pretend otherwise
3. **The decision you must make:** whether to keep the submission, reframe it around the energy/impossibility story, or stop here